

input token budget

detector (pooled = trained on all stages at once)

	M1	M2	M3	M4	M5	pooled
32	0.53	0.52	0.53	0.53	0.53	0.54
64	0.67	0.68	0.71	0.68	0.73	0.70
128	0.82	0.74	0.86	0.86	0.87	0.86
256	0.88	0.77	0.89	0.90	0.91	0.91
512	0.89	0.78	0.90	0.91	0.91	0.92
1024	0.90	0.78	0.90	0.91	0.91	0.92

plateau: within 0.01 of full length degraded: truncation costs real accuracy at chance: the task is not solvable

every cell is mean accuracy over the five test stages

full length is 1024 tokens, and 512 matches it - the plateau is not a single reference point

one paragraph is roughly 130-200 tokens, which puts the realistic case at the knee